

# FRIS — STEP 2: FEATURE INFERENCE CONTRACT (FROZEN)

**Status:** FINAL — this is the single source of truth for raw → engineered inference

This document defines the exact behavior, boundaries, and guarantees of the **feature inference layer**. No model logic. No explainability. No shortcuts.

---

## PURPOSE (ONE SENTENCE)

Convert **one raw transaction dict** into the **exact engineered feature DataFrame** used during training.

Nothing more.

---

## FILE OWNERSHIP

**File:** `src/feature_inference.py`

Why a separate file: - `features.py` defines transformations - `feature_inference.py` applies them in inference context - `models.py` must never know how features are created

This file is the **only legal entry point** for inference feature creation.

---

## PUBLIC API

```
def prepare_features(raw_input: dict) -> pd.DataFrame
```

This function is called by: - `models.predict()` - `models.predict_proba()` - `explain.explain_transaction()` - sanity / verification scripts

---

## INPUT CONTRACT

### raw\_input

- Type: `dict`
- Represents **one transaction only**
- Must contain the raw columns expected by training

Minimum required keys (current system): - Time - Amount

Notes: - No batch support - No list of dicts - No default filling here

---

## OUTPUT CONTRACT

- Type: pd.DataFrame
- Shape: (1, N)
- Columns: exactly the same as feature\_pipeline(..., fit=True) produced during training
- Order: not required, but column set must match

This DataFrame is **model-ready** but contains **no model outputs**.

---

## INTERNAL EXECUTION STEPS (MANDATORY ORDER)

### 1. Convert raw dict → single-row DataFrame

```
df_raw = pd.DataFrame([raw_input])
```

### 2. Load preprocessors from disk

3. scaler
4. encoders
5. PCA
6. must be the exact objects fit during training

### 7. Call feature pipeline in inference mode

```
df_eng, _ = feature_pipeline(  
    df_raw,  
    fit=False,  
    preprocessors=preprocessors  
)
```

### 8. Column contract assertion (HARD FAIL)

9. Load training\_feature\_columns.json
10. Compare against df\_eng.columns

11. Missing or extra columns → raise error

## 12. Return engineered DataFrame

---

### WHAT `prepare_features` MUST NOT DO

- No model loading
- No model inference
- No probability computation
- No anomaly scoring
- No latent extraction
- No clustering
- No SHAP
- No thresholding
- No try/except swallowing errors

If something breaks here, the system **must crash loudly**.

---

### WHAT `prepare_features` GUARANTEES

If this function returns successfully:

- Feature engineering is identical to training
- All downstream models can safely consume the output
- Explainability will operate on the same features as prediction

This is the **single choke point** for correctness.

---

## RELATIONSHIP TO OTHER FILES

### `features.py`

- Defines how features are created
- Stateless, reusable, inference-agnostic

### `feature_inference.py` (this file)

- Owns raw → engineered transition
- Loads state (preprocessors)
- Enforces training contract

### `models.py`

- Consumes engineered features
- Computes model signals
- Performs stacking and labeling

## **explain.py**

- Explains XGBoost using engineered features
  - Must never re-engineer features
- 

## **INFERENCE PIPELINE (FINAL, CANONICAL)**

```
raw_input (dict)
  ↓
prepare_features()
  ↓
engineered_df
  ↓
models.py (base models + stacker)
  ↓
score + label
```

Explainability uses the **same engineered\_df**.

---

## **STEP 2 OUTPUT (WHAT IS NOW FROZEN)**

- Feature inference responsibility is isolated
- Ownership boundaries are clean
- Raw → engineered path is unambiguous
- No file overlaps responsibility

## **READY FOR STEP 3**

Next step: **Refactor** `models.py` to consume `prepare_features()` **output and delete validate\_input()**.